


Markov Decision Process

Collect reward
Stochastic transitions



$\Pr(s'|s, a)$: Probability we transition to state s' if we choose action a in states s

MDP: $\Pr(s'|s, a)$, MP: $\Pr(s'|s)$



$R(s)$: Reward function

$\pi(s)$: Policy (Example)

$U([s_0, s_1, \dots, s_T])$: Utility of a sequence of states of fixed time

$$= \sum_{t=0}^T R(s_t)$$

$$= \sum_{t=0}^T \gamma^t R(s_t) \quad \gamma \in (0, 1]$$

(Decisions that help u now vs planning for later)

Value Iteration

This leads to $U^*(s) = \mathbb{E}_{\Pr(s_0, s_1, \dots, s_T) | s_0=s, \pi} \left[\sum_{t=0}^T \gamma^t R(s_t) \right]$ $G(\gamma)$:= (return, thing u are averaging)

Given all future trajectories given a s_0 and π , what is the average reward?

$\pi_s^* = \arg \max_{\pi} U^*(s) = \pi_s^*$ for any s' max gives you largest value

no one action, but probability of choosing an action

If policy stochastic

$$U^{\pi_s}(s) = \sum_a \pi_s(a) \sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma U^{\pi_s}(s')]$$

Bellman Equation: $U^*(s) = \max_a \left[\sum_{s'} P(s'|s, a) (R(s, a, s') + \gamma U^*(s')) \right]$

Recursive def (how it is solved in practice)

From U^* u can get π^* , but not vice versa

Initial Guesses for π_0 for $\forall s_0 \in S$

$U^{\pi_0}(s) = \sum_{s'} P(s'|s, \pi_0(s)) (R(s, \pi_0(s), s') + \gamma U^{\pi_0}(s'))$

$\pi_{t+1}(s) = \arg \max_a \sum_{s'} P(s'|s, a) (R(s, a, s') + \gamma U^{\pi_t}(s'))$

Value Iteration
u can say $s' = s_{t+1}$ & $s = s_t$ here

- Computes utility for every state
- Needs exact transition model
- Needs fully observe state
- Needs to know exact reward for each state

Partially-observable MDPs

$(S, A, T, R, \Omega, O, \gamma)$ a_t = action at time t

Hidden $\left\{ \begin{array}{l} \text{States } S \\ \text{Actions } A \end{array} \right.$ T : Transition fct: $T(s'|s, a) = P(s_{t+1}=s' | s_t=s, a_t=a)$
 R : Reward fct: $R(s, a, s')$

Observable $\left\{ \begin{array}{l} \text{Observations } \Omega: \text{Sensor readings and tokens the agent can actually perceive} \\ \text{Observation fct } O: O(o|s, a) = P(o_{t+1}=o | s_{t+1}=s', a_t=a) \\ \text{Discount factor } \gamma \end{array} \right.$

agent doesn't know exact state s , so maintains belief state $b(s)$

$$\sum_{s \in S} b(s) = 1$$

Belief updates: $b'(s') = \gamma \cdot O(o|s', a) \sum_{s \in S} T(s'|s, a) b(s)$ normalization $\gamma = \frac{P(o|b, a)}{P(o|b, a)}$

(Bayes rule)

Now, we do $\pi(b(s)) \rightarrow$ action a instead of $\pi(s) \rightarrow a$

1. you have b_t
2. Then you have $a_t = \pi_t(b_t)$
3. O_{t+1}
4. $b_{t+1}(s') = \gamma O(o_{t+1}|s', a_t) \sum_{s \in S} P(s'|s, a_t) b_t(s)$
5. $U^{\pi_t}(b_t) = R(b_t, \pi_t(b_t)) + \gamma \sum_{a_{t+1}} P(o_{t+1}|b_t, \pi_t(b_t)) U^{\pi_t}(b_{t+1})$
6. $\pi_{t+1}(b) = \arg \max_a [R(b_t, a) + \gamma \sum_{o_{t+1}} P(o_{t+1}|b_t, a) U^{\pi_t}(b_{t+1})]$

Policy Gradient Algorithm & PPO

$$J = \mathbb{E}_{\tau \sim \pi_{\theta}} [R(\tau)] = \int \pi_{\theta}(\tau) R(\tau) d\tau \quad \mathbb{E}[X] = \int x P(x)$$

$$\nabla_{\theta} J = \int R(\tau) [\pi_{\theta} \cdot \nabla_{\theta} \log \pi_{\theta}(\tau)] d\tau = \mathbb{E}_{\tau \sim \pi_{\theta}} (A(\tau) \nabla_{\theta} \log \pi_{\theta}(\tau))$$

- Reinforce Alg:
1. Reinforce one full trajectory using π_{θ}
 2. Calc $R(\tau)$ from black box τ = trajectory/complete response
 3. Calc score: $\nabla_{\theta} \log \pi_{\theta}(\tau)$
 4. $\theta \leftarrow \theta + \alpha \cdot [R(\tau) \nabla_{\theta} \log \pi_{\theta}(\tau)]$

Now say: $A(s_t, a_t) = R_t - V(s_t)$

$$\nabla_{\theta} J = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) A(s_t, a_t) \right]$$


Instability tho, i.e massive gradient steps, so we need "safety harness"

PPO

$$r_t(\theta) = \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_{old}}(a_t | s_t)} \quad L^{clip}(\theta) = \mathbb{E}_t [\min(r_t(\theta) A_t, \text{clip}(r_t(\theta), 1-\epsilon, 1+\epsilon) A_t)]$$


1. $A_t > 0$ $\min(r_t(\theta) A_t, (1+\epsilon) A_t)$ (Ceiling)
- if $r_t > 1+\epsilon$ objective goes flat (floor)
2. $A_t < 0$; $r_t < 1-\epsilon$

Actor-Critic Model



Value-based Alg: Learn value fct $Q(s, a)$, Implicit π ; $\arg \max Q$ (Challenges in continuous action space)

Policy-based alg: No value fct learn $\pi_{\theta}(a|s)$ or $\mu_{\theta}(a|s)$ (High variance)



Actor-critic model (best of both worlds)

For recursive stuff they guess, use A assuming B is right, then fix B , etc

TD: A update, B update $\rightarrow A$ update; $C \rightarrow B \rightarrow A$


Monte-Carlo: Wait til ending, then go backwards once

(only if end episode exists/finite)

Actor: Decides which action to take (Policy-based method)

Critic: How good was action & how to adjust (Value-based)

$$Q(s_t, a_t) = \mathbb{E} \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k} \right]$$



Obtain data experience $[s_t, a_t, r_t, s_{t+1}]$

Online Training: One experience at a time

Offline Training: In batches

Training

(4) Compute $Q_w(s_t, a_t)$

(5) $\Delta \theta = \alpha \nabla_{\theta} \log \pi_{\theta}(s_t, a_t) Q_w(s_t, a_t)$ Actor updates policy params θ using Q-val $\rightarrow$ policy grad

(6) $\Delta w = \beta [R(s_t, a_t) + \gamma Q_w(s_{t+1}, a_{t+1}) - Q_w(s_t, a_t)] \nabla_w Q_w(s_t, a_t)$

(only if end episode exists/finite)

δ CTD error

Update Critic: $(R_t - V(s_t))^2$

Update Actor: $\nabla_{\theta} \log \pi_{\theta} - A$

Advantage: $A = R - V$

Predicted Value $V(s)_t$

Critic V_{ϕ}

Env/Rewards

action a_t

Actor π_{θ}

State s_t

R_t

s_{t+1}

R_t

R_t

R_t

R_t

R_t

R_t

R_t

R_t

R_t

R_t

R_t

R_t

R_t

R_t

R_t

R_t

R_t

R_t

Problem: Instability